


# Introduction to Computer Networking

Guy Leduc

## Chapter 2 Application Layer



*Computer Networking: A  
Top Down Approach,  
8<sup>th</sup> edition.*

Jim Kurose, Keith Ross  
Pearson, 2020

# Chapter 2: outline

2.1 Principles of network applications

2.2 Web and HTTP

2.3 The Domain Name System (DNS)

2.4 Socket programming with UDP and TCP

# Application Layer: Overview

## Our goals:

- ❑ Conceptual implementation aspects of application-layer protocols
  - ❖ transport-layer service models
  - ❖ client-server paradigm
  - ❖ peer-to-peer paradigm
- ❑ learn about protocols by examining popular application-layer protocols
  - ❖ HTTP
  - ❖ DNS
- ❑ programming network applications
  - ❖ socket API

## Some network apps

- ❑ social networking
- ❑ web
- ❑ text messaging
- ❑ e-mail
- ❑ multi-user network games
- ❑ streaming stored video  
(YouTube, Netflix, ...)
- ❑ P2P file sharing
- ❑ voice over IP (e.g., Skype)
- ❑ real-time video  
conferencing
- ❑ Internet search
- ❑ remote login
- ❑ ...

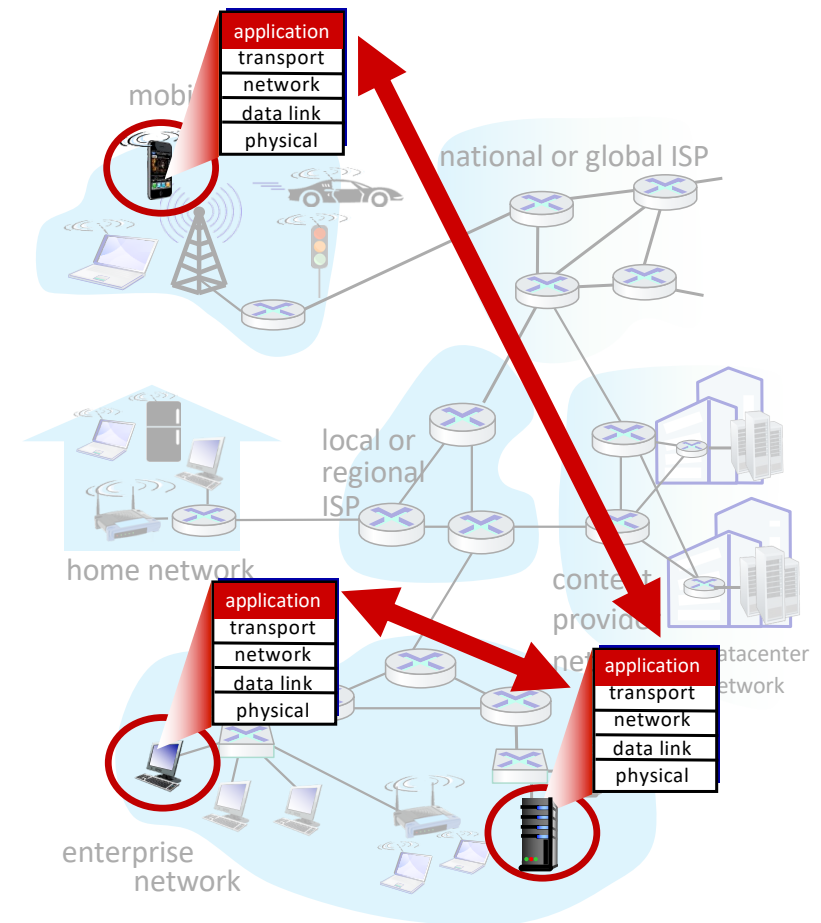
# Creating a network app

write programs that:

- run on (different) end systems
- communicate over network
- e.g., web server software communicates with browser software

no need to write software for network-core devices

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation



# Application architectures

## Possible structure of applications:

- ❖ Client-server
- ❖ Peer-to-peer (P2P)


# Client-server paradigm

## server:

- always-on host
- permanent IP address
- often in data centers, for scaling

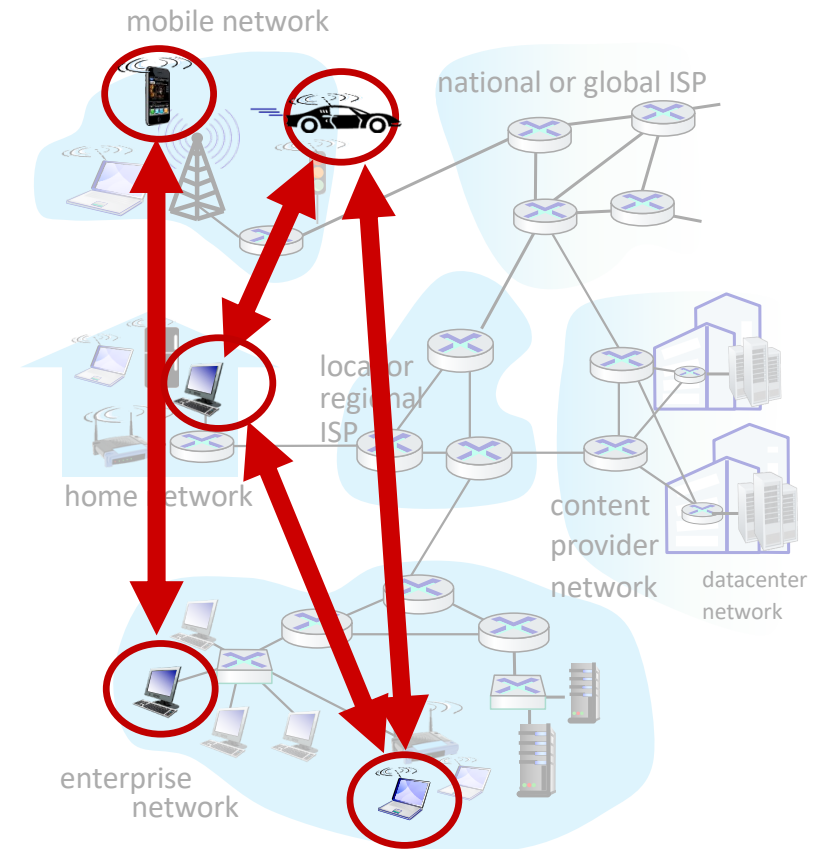
## clients:

- contact, communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do *not* communicate directly with each other
- examples: HTTP, IMAP, FTP



# Peer-peer (P2P) architecture

- *no* always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
  - *self scalability* – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
  - complex management
- example: P2P file sharing





# Hybrid of client-server and P2P

## Skype

- ❖ voice-over-IP P2P application
- ❖ centralized server: finding address of remote party
- ❖ client-client connection: direct (not through server)

## Text messaging

- ❖ chatting between two users is P2P
- ❖ centralized service: client presence detection/location

## BitTorrent

- ❖ exchanging file chunks between users is P2P
- ❖ tracker: maintains list of peers participating in torrent

# Processes communicating

*process*: program running within a host

- within same host, two processes communicate using *inter-process communication* (defined by OS)
- processes in different hosts communicate by exchanging *messages*

clients, servers


*client process*: process that initiates communication

*server process*: process that waits to be contacted

- ❖ aside: applications with P2P architectures have client processes & server processes

# Sockets

- ❑ process sends/receives messages to/from its **socket**
- ❑ socket analogous to door
  - ❖ sending process shoves message out door
  - ❖ sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process
  - ❖ two sockets involved: one on each side



## Addressing processes

- ❑ to receive messages, process must have *identifier*
- ❑ host device has (at least) one IP address
- ❑ Q: does IP address of host on which process runs suffice for identifying the process?
  - A: no, *many* processes can be running on same host
- ❑ *identifier* includes both **IP address** and **port number** associated with process on host
- ❑ example port numbers:
  - ❖ HTTP server: 80
  - ❖ mail server: 25
- ❑ to send HTTP message to gaia.cs.umass.edu web server:
  - ❖ **IP address:** 128.119.245.12
  - ❖ **port number:** 80
- ❑ more shortly...

# An application-layer protocol defines

- ❑ **Types of messages exchanged:**

- ❖ e.g., request, response

- ❑ **Message syntax:**

- ❖ what fields in messages & how fields are delineated

- ❑ **Message semantics:**

- ❖ meaning of information in fields

- ❑ **Rules** for when and how processes send & respond to messages

## **Open protocols:**

- ❑ defined in RFCs, everyone has access to protocol definition
- ❑ allows for interoperability
- ❑ e.g., HTTP, SMTP

## **Proprietary protocols:**

- ❑ e.g., Skype

# What transport service does an application need?

## Data integrity

- ❑ some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- ❑ other apps (e.g., audio) can tolerate some loss

## Timing

- ❑ some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

## Throughput

- ❑ some apps (e.g., multimedia) require minimum amount of throughput to be “effective”
- ❑ other apps (“elastic apps”) make use of whatever throughput they get

## Security

- ❑ encryption, data integrity, ...

## Transport service requirements of common applications

Application	Data loss	Throughput	Time Sensitive?
file transfer/download	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video:10kbps-5Mbps	yes, 10's msec
streaming audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	few kbps up	yes, 10's msec
text messaging	no loss	elastic	yes and no

# Internet transport layer services

## TCP service:

- ❑ *reliable transport* between sending and receiving process
- ❑ *flow control*: sender won't overwhelm receiver
- ❑ *congestion control*: throttle sender when network overloaded
- ❑ *does not provide*: timing, minimum throughput guarantees, security
- ❑ *connection-oriented*: setup required between client and server processes

useful for e.g. file transfer, email, web, streaming stored video

## UDP service:

- ❑ *unreliable data transfer* between sending and receiving process
- ❑ *does not provide*: reliability, flow control, congestion control, timing, throughput guarantee, security, connection setup

Q: why is there a UDP?

useful e.g. for internet telephony, interactive games, DNS



# Securing TCP

## Vanilla TCP & UDP sockets

- ❑ no encryption
- ❑ cleartext (e.g. passwords) sent into socket traverse Internet in cleartext (!)

## Transport Layer Security (TLS)

- ❑ provides encrypted TCP connections
- ❑ data integrity
- ❑ end-point authentication

## TLS implemented in app layer

- ❑ between apps and TCP
- ❑ apps use TLS libraries, that use TCP in turn

## TLS socket API

- ❑ cleartext sent into socket traverse Internet *encrypted*

More on TLS in chapter 8 of book

# Chapter 2: outline

2.1 Principles of network applications

2.2 Web and HTTP

2.3 The Domain Name System (DNS)

2.4 Socket programming with UDP and TCP

# Web and HTTP

First, a quick review...

- ❑ Web page consists of **objects**, each of which can be stored on different Web servers
- ❑ Object can be HTML file, JPEG image, Java applet, audio file,...
- ❑ Web page consists of **base HTML-file** which includes **several referenced objects**, each addressable by a **URL** (Universal Resource Locator), e.g.:

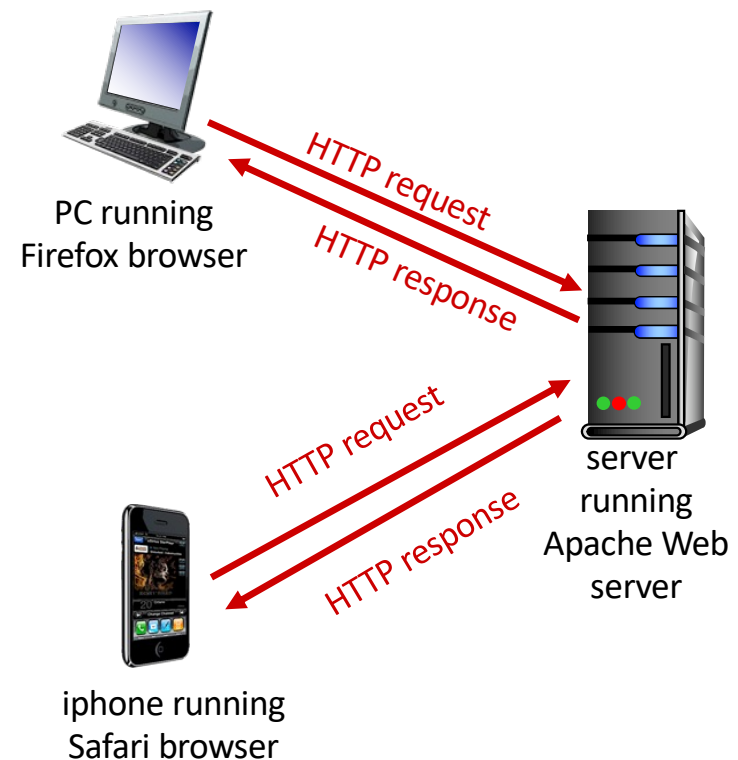
`http://www.someschool.edu:8080/someDept/pic.gif`

protocol name      host name or IP addr      port (if non standard)      path name

# HTTP overview

## HTTP: HyperText Transfer Protocol

- Web's application layer protocol
- client/server model
  - ❖ *client*: browser that requests, receives, (using HTTP protocol) and “displays” Web objects
  - ❖ *server*: Web server sends (using HTTP protocol) objects in response to requests



## HTTP overview (continued)

### HTTP uses TCP (why?):

- ❑ client initiates TCP connection (creates socket) to server, port 80
- ❑ server accepts TCP connection from client
- ❑ HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- ❑ TCP connection closed

### HTTP is “stateless”

- ❑ server maintains no information about past client requests

aside  
Protocols that maintain “state”  
are complex!

- ❑ past history (state) must be maintained
- ❑ if server/client crashes, their views of “state” may be inconsistent, must be reconciled

## HTTP connections: two types

### *Non-persistent HTTP*

1. TCP connection opened
2. at most one object sent over TCP connection
3. TCP connection closed

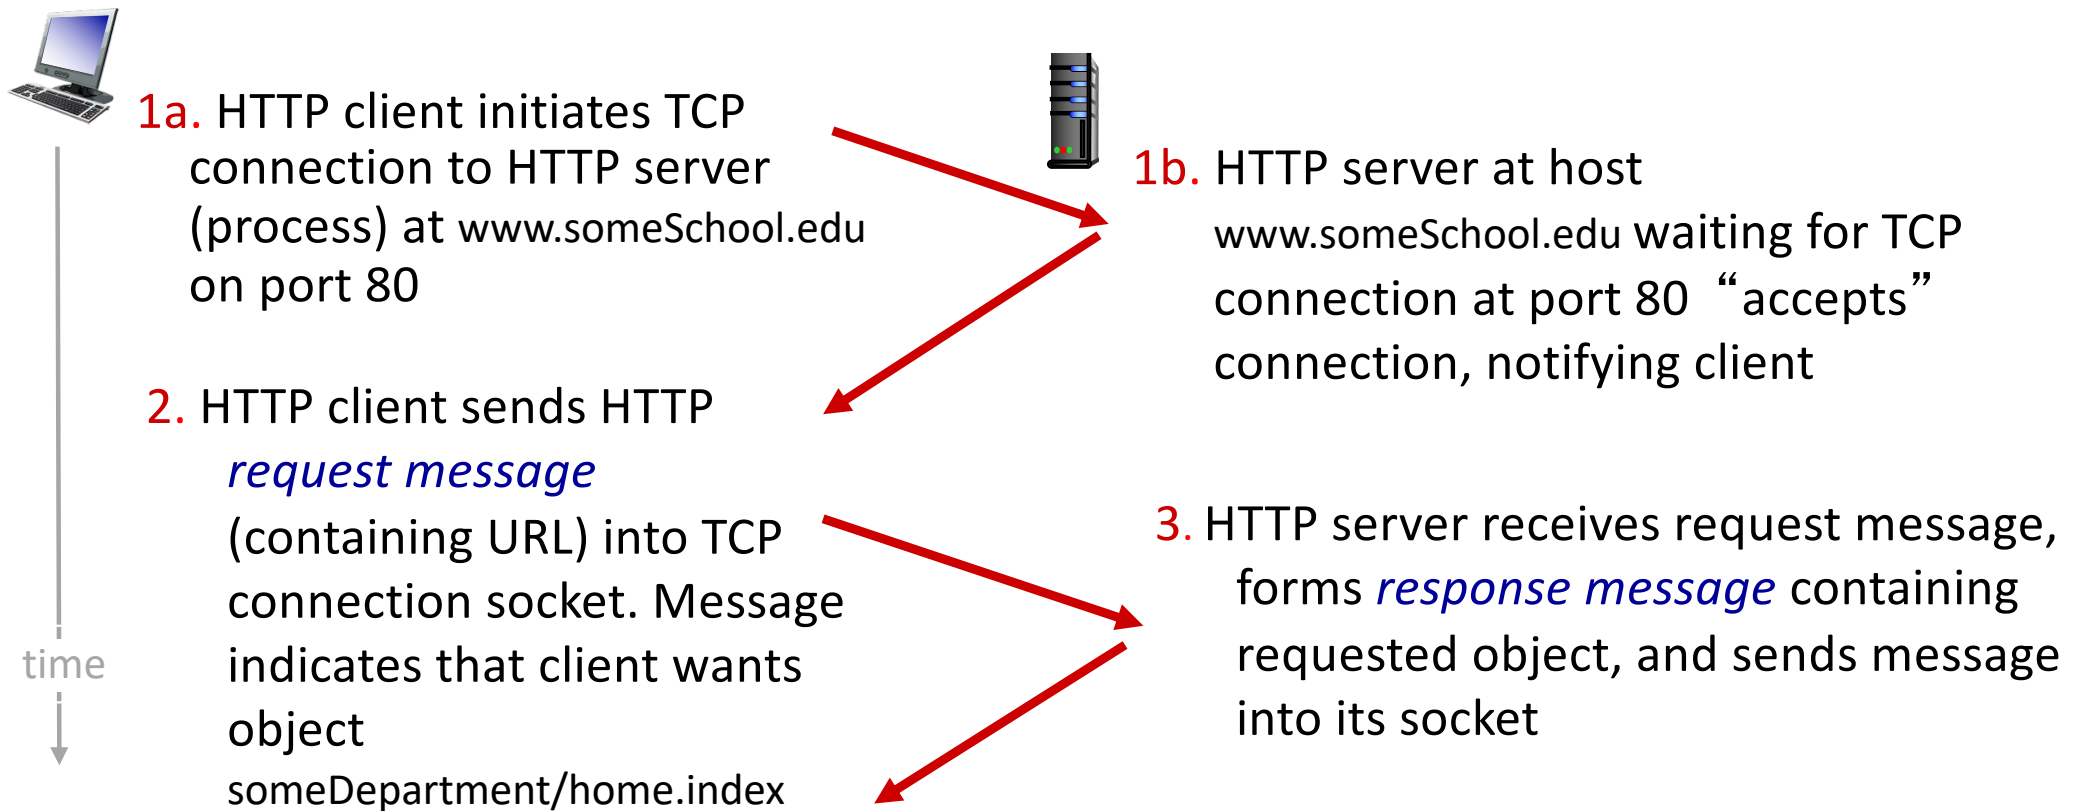
downloading multiple objects required multiple connections

### *Persistent HTTP*

- TCP connection opened to a server
- multiple objects can be sent over *single* TCP connection between client, and that server
- TCP connection closed

# Non-persistent HTTP: example

User enters URL: `www.someSchool.edu/someDepartment/home.index`  
(containing text, references to 10 jpeg images)



## Non-persistent HTTP: example (cont.)

User enters URL: `www.someSchool.edu/someDepartment/home.index`  
(containing text, references to 10 jpeg images)



4. HTTP server closes TCP connection.

5. HTTP client receives response message containing html file, displays html. Parsing html file, finds 10 referenced jpeg objects

6. Steps 1-5 repeated for each of 10 jpeg objects

time



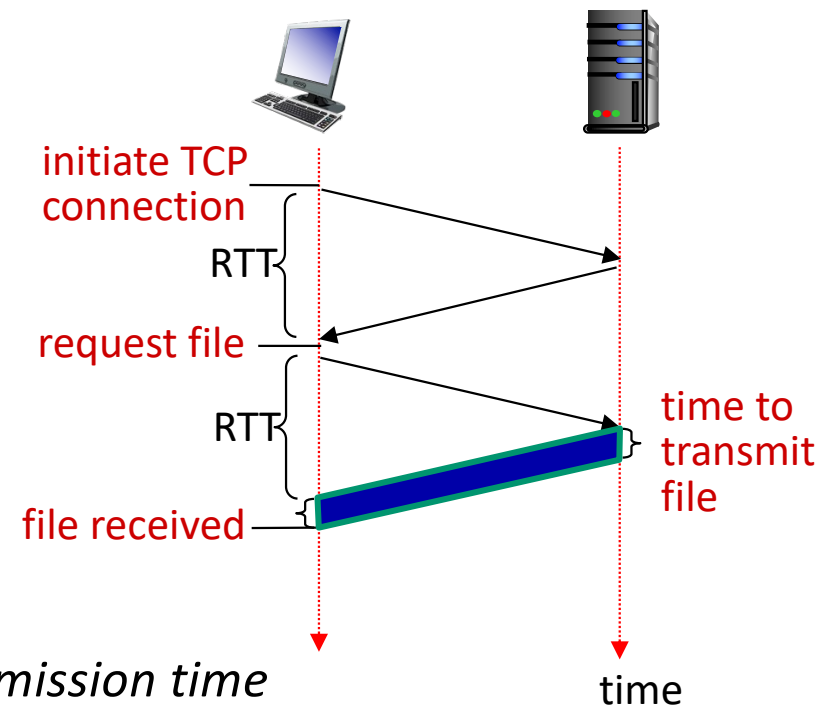


# Non-persistent HTTP: response time

**RTT (definition):** Round Trip Time :  
time for a small packet to travel from  
client to server and back

**HTTP response time (per object):**

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- object/file transmission time



*Non-persistent HTTP response time =  $2RTT + \text{file transmission time}$*

File transmission time = file size divided by the average throughput of underlying TCP connection

## Persistent HTTP (HTTP 1.1)

### Non-persistent HTTP issues:

- ❑ requires 2 RTTs per object
- ❑ OS overhead for *each* TCP connection
- ❑ browsers often open parallel TCP connections to fetch referenced objects in parallel

### Persistent HTTP (HTTP1.1):

- ❑ server leaves connection open after sending response
- ❑ subsequent HTTP messages between same client/server sent over open connection
- ❑ client sends requests as soon as it encounters a referenced object
- ❑ as little as one RTT for all the referenced objects when requests are pipelined (cutting response time in half)

# HTTP request message

□ two types of HTTP messages: *request, response*

□ **HTTP request message:**

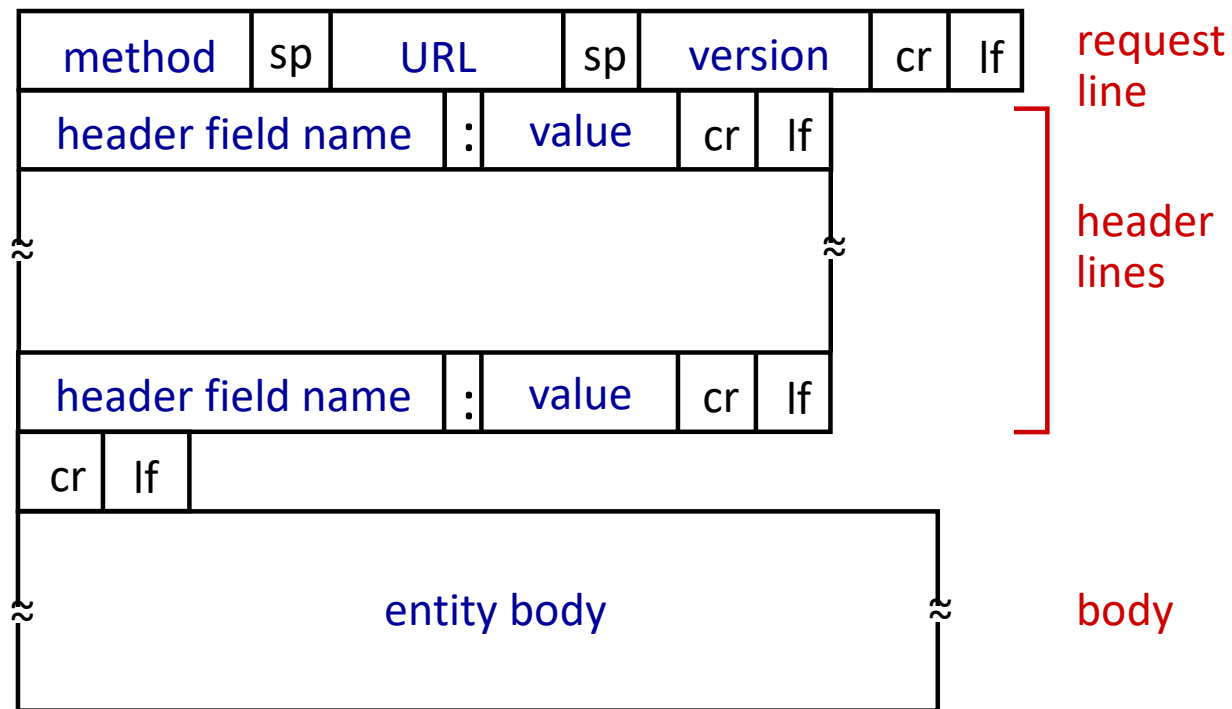
❖ ASCII (human-readable format)

The diagram illustrates the structure of an HTTP request message. It shows the following components and their corresponding annotations:

- request line (GET, POST, HEAD commands):** Points to the first line of the message: `GET /index.html HTTP/1.1\r\n`. A separate annotation points to the `\r` and `\n` characters, identifying them as the **carriage return character** and **line-feed character** respectively.
- header lines:** A bracket groups the subsequent lines: `Host: www-net.cs.umass.edu\r\n`, `User-Agent: Firefox/3.6.10\r\n`, `Accept: text/html,application/xhtml+xml\r\n`, `Accept-Language: en-us,en;q=0.5\r\n`, `Accept-Encoding: gzip,deflate\r\n`, `Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n`, `Keep-Alive: 115\r\n`, and `Connection: keep-alive\r\n`. An annotation points to the `\r\n` at the end of the first header line, stating: **carriage return, line feed at start of line indicates end of header lines**.
- Persistent mode:** A bracket groups the last two header lines, with an annotation stating: **Persistent mode**.
- Why needed?:** An annotation points to the `Host` header line.

```
GET /index.html HTTP/1.1\r\nHost: www-net.cs.umass.edu\r\nUser-Agent: Firefox/3.6.10\r\nAccept: text/html,application/xhtml+xml\r\nAccept-Language: en-us,en;q=0.5\r\nAccept-Encoding: gzip,deflate\r\nAccept-Charset: ISO-8859-1,utf-8;q=0.7\r\nKeep-Alive: 115\r\nConnection: keep-alive\r\n\r\n
```

# HTTP request message: general format



# Other HTTP request messages

## POST method:

- web page often includes form input
- user input sent from client to server in entity body of HTTP POST request message
- more generally POST sends data to a remote object specified in URL field

GET method (another technique for sending data to server):

- include user data in URL field of HTTP GET request message (following a '?'):

`www.somesite.com/animalsearch?monkeys&banana`

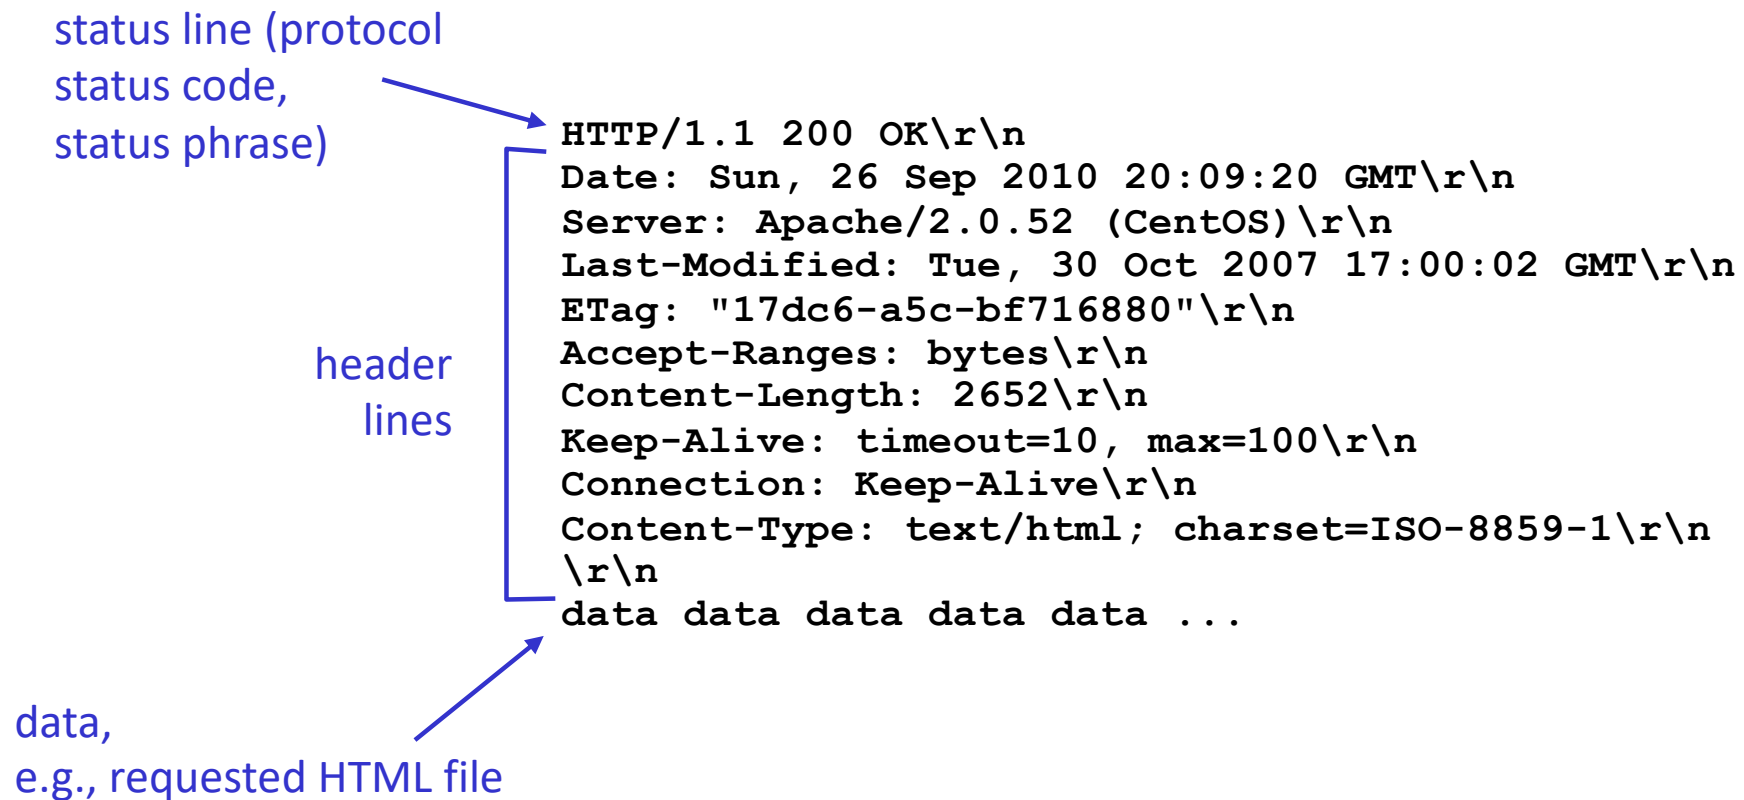
## HEAD method:

- requests headers (only) that would be returned *if* specified URL were requested with an HTTP GET method

## PUT method:

- uploads new file (object) to server
- completely replaces file/object that exists at specified URL with content in entity body of PUT HTTP request message

# HTTP response message



## HTTP response status codes

- status code appears in 1<sup>st</sup> line in server-to-client response message

- some sample codes:

200 OK

- ❖ request succeeded, requested object later in this msg

301 Moved Permanently

- ❖ requested object moved, new location specified later in this msg (in Location: field)

400 Bad Request

- ❖ request msg not understood by server

404 Not Found

- ❖ requested document not found on this server

505 HTTP Version Not Supported

# Trying out HTTP (client side) for yourself

## 1. Telnet to your favorite Web server:

**telnet gaia.cs.umass.edu 80**

- opens TCP connection to port 80 (default HTTP server port) at gaia.cs.umass.edu.
- anything typed in will be sent to port 80 at gaia.cs.umass.edu

## 2. type in a GET HTTP request:

**GET /kurose\_ross/interactive/index.php HTTP/1.1**

**Host: gaia.cs.umass.edu**

- by typing this in (hit carriage return twice), you send this minimal (but complete) GET request to HTTP server

## 3. look at response message sent by HTTP server!

(or use Wireshark to look at captured HTTP request/response)



## Maintaining user/server state: cookies

Recall: HTTP GET/response interaction is *stateless*

- no notion of multi-step exchanges of HTTP messages to complete a Web “transaction”
- no need for client/server to track “state” of multi-step exchange
- all HTTP requests are independent of each other

# Maintaining user/server state: cookies

Web sites and client browser use *cookies* to maintain some state between transactions

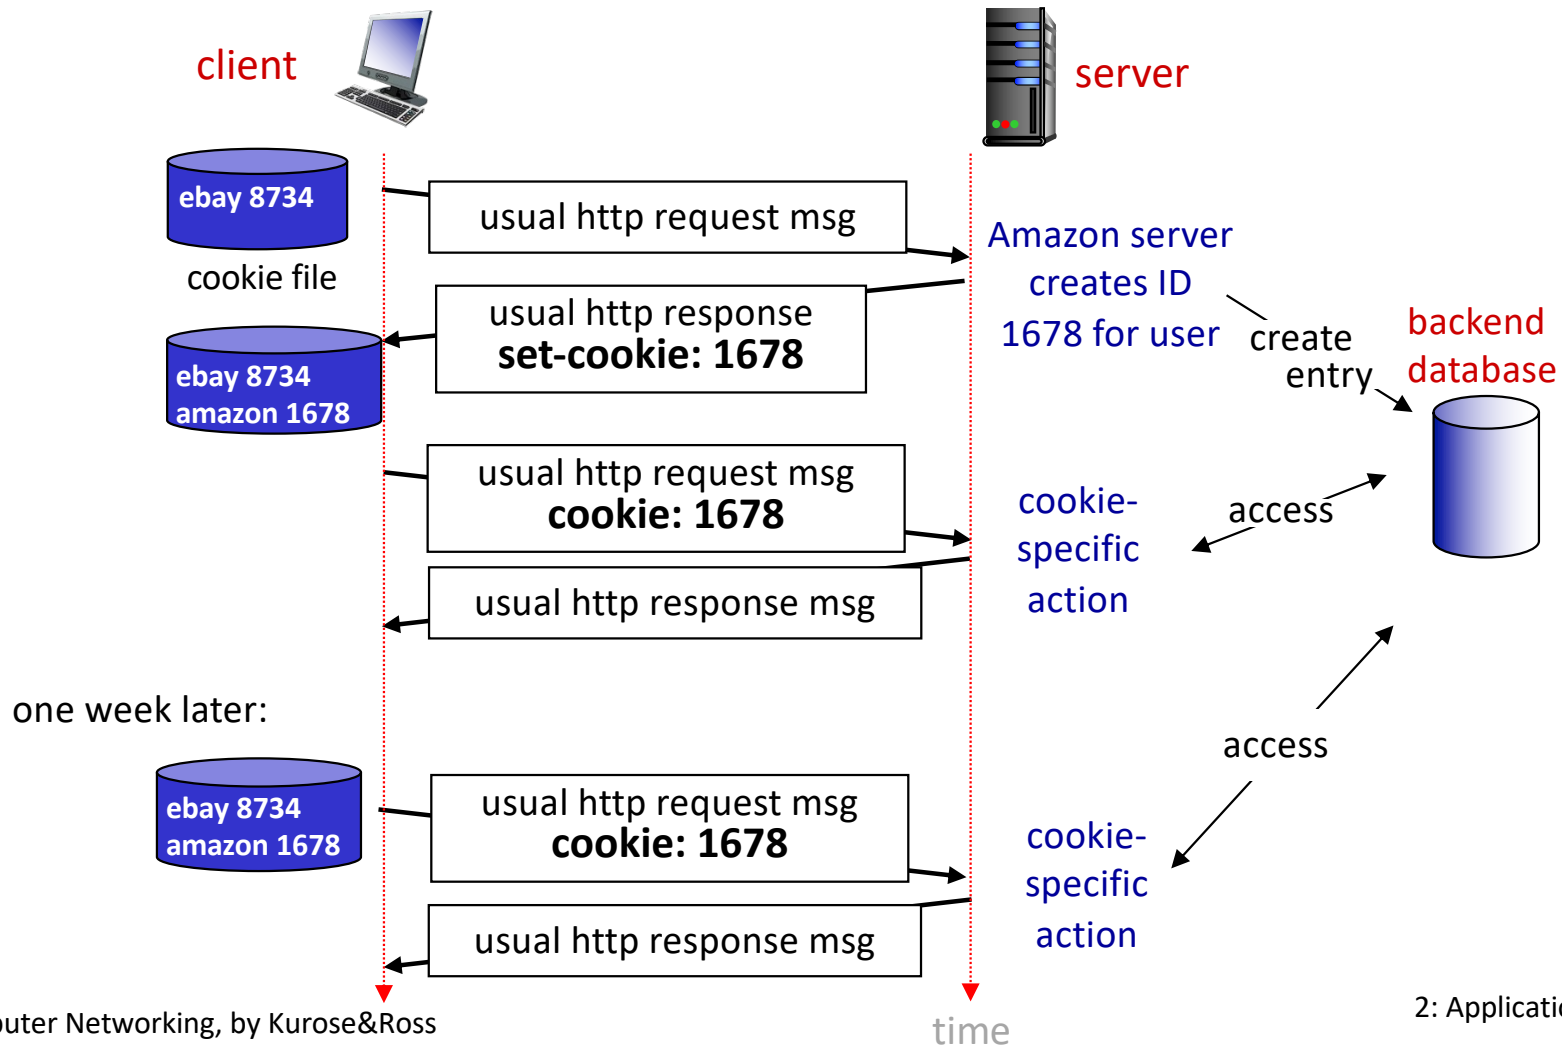
## *four components:*

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

## Example:

- Susan uses browser on laptop, visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
  - unique ID (aka “cookie”)
  - entry in backend database for ID
- subsequent HTTP requests from Susan to this site will contain cookie ID value, allowing site to “identify” Susan

## Maintaining user/server state: cookies



# HTTP cookies: comments

## *What cookies can be used for:*

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

## *Challenge: How to keep state:*

- protocol endpoints: maintain state at sender/receiver over multiple transactions
- cookies: HTTP messages carry state
- any security issue?

aside


*cookies and privacy:*

- cookies permit sites to *learn* a lot about you on their site.
- third party persistent cookies (tracking cookies) allow common identity (cookie value) to be tracked across multiple web sites

## Web caches (proxy server)

*goal:* satisfy client request without involving origin server

- ❑ user configures browser to point to a **Web cache**
- ❑ browser sends all HTTP requests to cache
  - ❖ *if* object in cache: cache returns object to client
  - ❖ *else* cache requests object from origin server, caches received object, then returns object to client



# Web caches (proxy servers)

- ❑ cache acts as both client and server
  - ❖ server for original requesting client
  - ❖ client to origin server
- ❑ typically cache is installed by ISP (university, company, residential ISP)

## Why Web caching?

1. reduce response time for client request: cache is closer to client
  2. reduce traffic on an institution's access link
  3. reduce load on servers
- 
- ❑ Internet is dense with caches: enables “poor” content providers to more effectively deliver content

# Caching example


## *Scenario:*

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- Web object size: 100K bits
- Average request rate from browsers to origin servers: 15/sec
  - average data rate to browsers: 1.50 Mbps

## *Performance:*

- LAN utilization: .0015
- access link utilization = .97
- end-end delay = Internet delay +  
access link delay + LAN delay  
= 2 sec + minutes +  $\mu$ secs

*problem: large  
delays at high  
utilization!*



# Caching example: buy a faster access link


## *Scenario:*

- access link rate: ~~1.54 Mbps~~ <sup>154 Mbps</sup>
- RTT from institutional router to server: 2 sec
- Web object size: 100K bits
- Avg request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

## *Performance:*

- LAN utilization: .0015
- access link utilization = ~~.97~~ <sup>.0097</sup>
- end-end delay = Internet delay +  
access link delay + LAN delay  
= 2 sec + ~~minutes~~ <sup>msecs</sup>

*Cost:* faster access link (expensive!) <sup>msecs</sup>





# Caching example: install a web cache


## *Scenario:*

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- Web object size: 100K bits
- Avg request rate from browsers to origin servers: 15/sec
  - avg data rate to browsers: 1.50 Mbps

## *Performance:*

- LAN utilization: .?
  - access link utilization = ?
  - average end-end delay = ?
- How to compute link utilization, delay?*


*Cost:* web cache (cheap!)



# Caching example: install a web cache

## Calculating access link utilization, end-end delay with cache:

- suppose cache hit rate is 0.4: 40% requests satisfied at cache, 60% requests satisfied at origin
- access link: 60% of requests use access link
- data rate to browsers over access link  
 $= 0.6 * 1.50 \text{ Mbps} = 0.9 \text{ Mbps}$
- utilization  $= 0.9 / 1.54 = 0.58$
- average end-end delay  
 $= 0.6 * (\text{delay from origin servers})$   
 $+ 0.4 * (\text{delay when satisfied at cache})$   
 $= 0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$



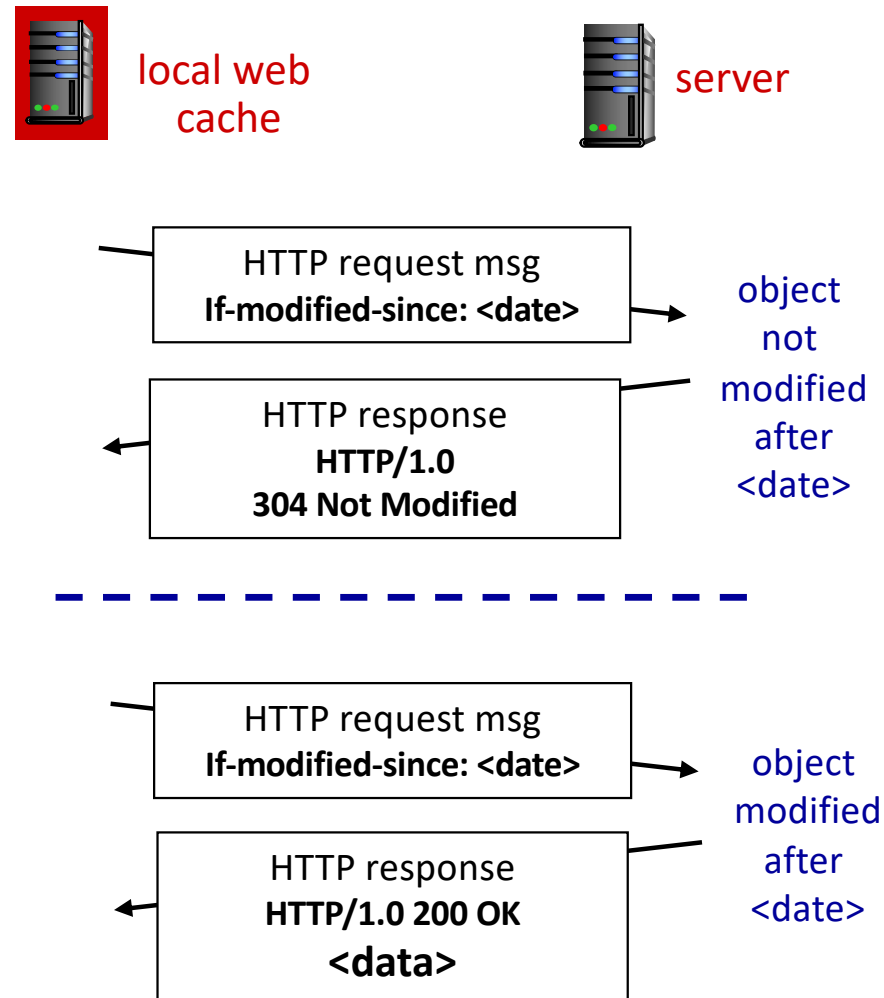
*lower average end-end delay than with 154 Mbps link (and cheaper too!)*

# Conditional GET

- ❑ **Goal:** don't send object if cache has up-to-date cached version
  - ❖ no object transmission delay
  - ❖ lower link utilization
- ❑ **cache:** specify date of cached copy in HTTP request

**If-modified-since: <date>**
- ❑ **server:** response contains no object if cached copy is up-to-date:

**HTTP/1.0 304 Not Modified**



# HTTP/2

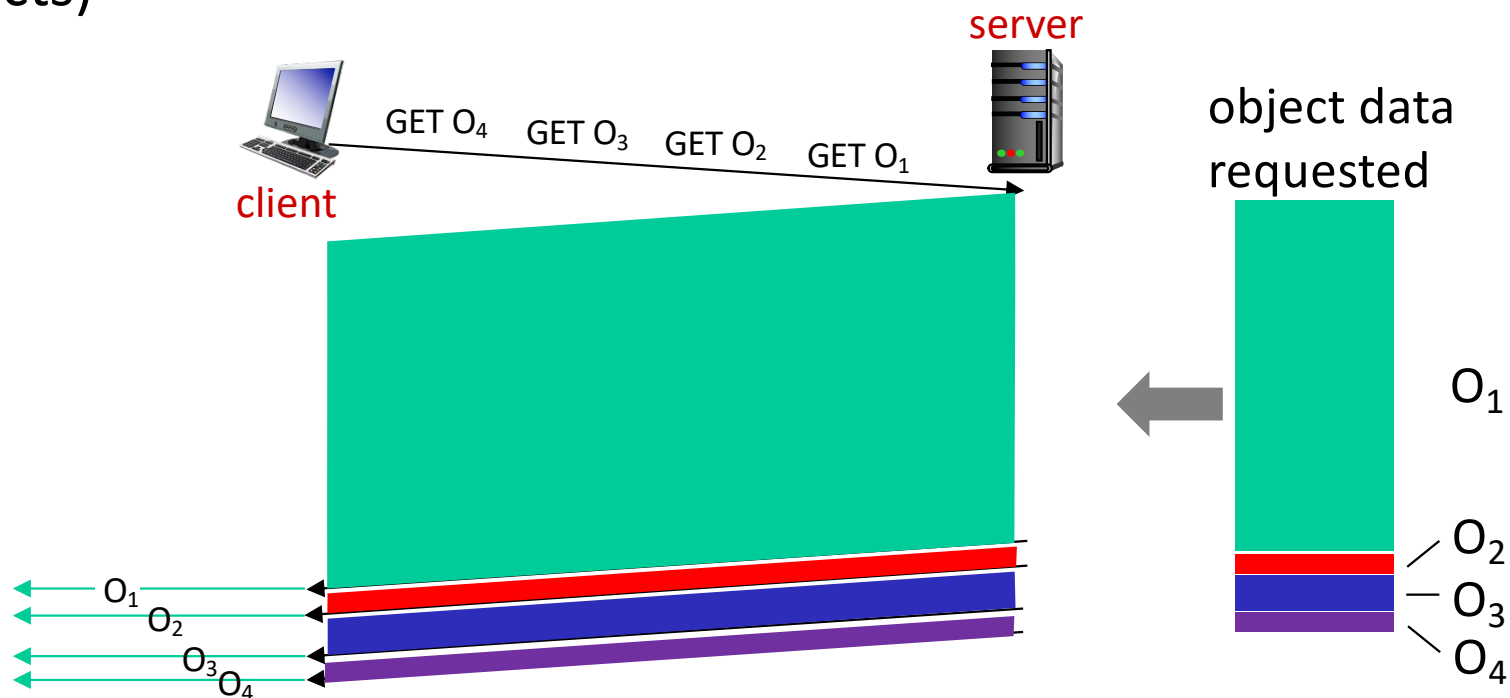
*Key goal:* decreased delay in multi-object HTTP requests

HTTP1.1: introduced **multiple, pipelined GETs** over single TCP connection

- server responds *in-order* (FCFS: first-come-first-served scheduling) to GET requests
- with FCFS, small object may have to wait for transmission (**head-of-line (HOL) blocking**) behind large object(s)
- loss recovery (retransmitting lost TCP segments) stalls object transmission

# HTTP1.1: HOL blocking

HTTP 1.1: client requests 1 large object (e.g., video file, and 3 smaller objects)



*objects delivered in order requested:  $O_2$ ,  $O_3$ ,  $O_4$  wait behind  $O_1$*

# HTTP/2

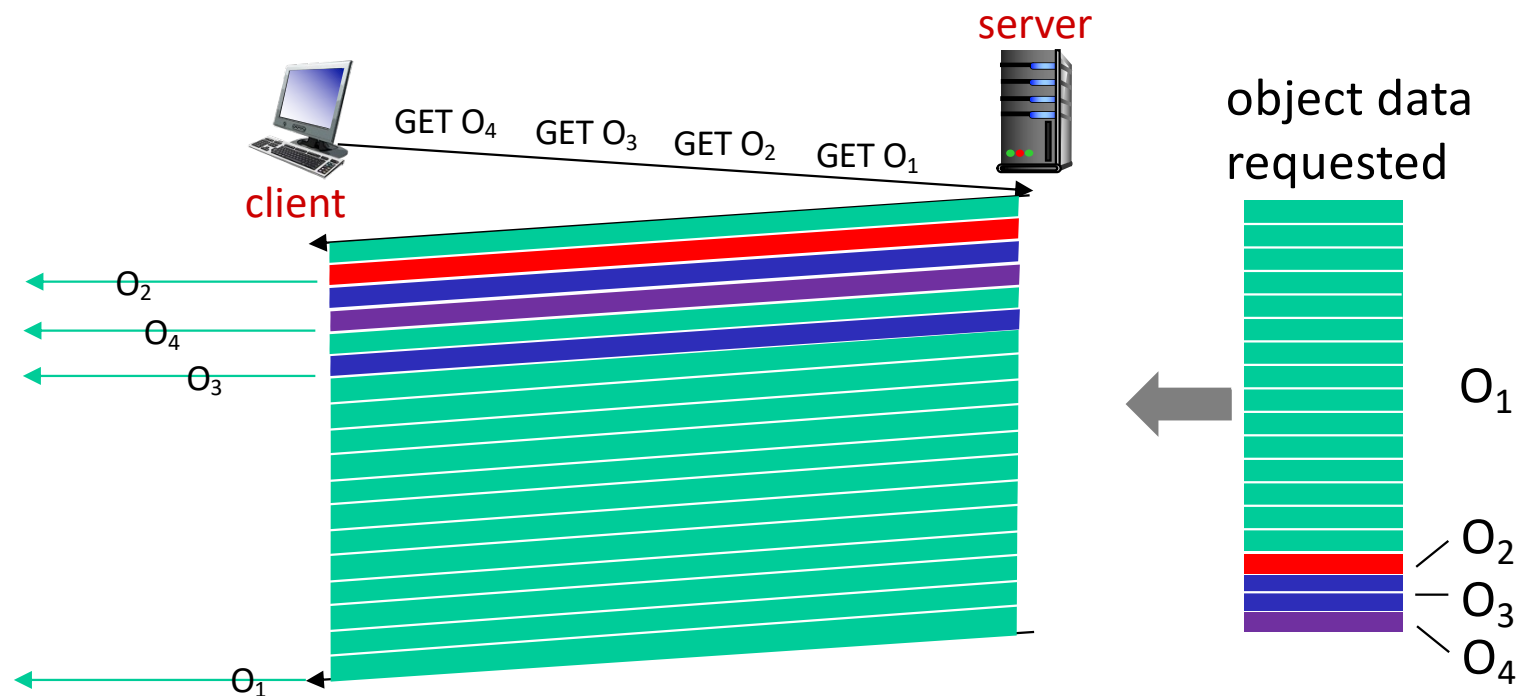
*Key goal:* decreased delay in multi-object HTTP requests

HTTP/2: [RFC 7540, 2015] increased flexibility at *server* in sending objects to client:

- methods, status codes, most header fields unchanged from HTTP 1.1
- transmission order of requested objects based on client-specified object priority (not necessarily FCFS)
- *push* unrequested objects to client
- divide objects into frames, schedule frames to mitigate HOL blocking

# HTTP/2: mitigating HOL blocking

HTTP/2: objects divided into frames, frame transmission interleaved



*O<sub>2</sub>, O<sub>3</sub>, O<sub>4</sub> delivered quickly, O<sub>1</sub> slightly delayed*

## HTTP/2 to HTTP/3

*Key goal:* decreased delay in multi-object HTTP requests

HTTP/2 over single TCP connection means:

- recovery from packet loss still stalls all object transmissions
  - as in HTTP 1.1, browsers have incentive to open multiple parallel TCP connections to reduce stalling, increase overall throughput
- no security over vanilla TCP connection
- **HTTP/3:** adds security, per object error- and congestion-control (more pipelining) over UDP
  - more on HTTP/3 in transport layer



## Chapter 2: outline

2.1 Principles of network applications

2.2 Web and HTTP

2.3 The Domain Name System (DNS)

2.4 Socket programming with UDP and TCP

# DNS: Domain Name System

**People:** many identifiers:

- ❖ SSN, name, passport #

**Internet hosts, routers:**

- ❖ IP address (32 bit) - used for addressing datagrams
- ❖ “name”, e.g., www.amazon.com used by humans

**Q:** how to map between IP address and name, and vice versa?

**Domain Name System:**

- ❑ *distributed database* implemented in hierarchy of many *name servers*
- ❑ *application-layer protocol*  
host, name servers communicate to *resolve* names (address/name translation)
  - ❖ note: core Internet function, implemented as application-layer protocol
  - ❖ complexity at network’s “edge”

# DNS: services, structure

## DNS services

- hostname to IP address translation
- host aliasing
  - canonical, alias names
- mail server aliasing
- load distribution
  - replicated Web servers: many IP addresses correspond to one name


## *Q: Why not centralize DNS?*

- single point of failure
- traffic volume
- distant centralized database
- maintenance

## *A: doesn't scale!*

- Comcast DNS servers alone: 600 billion DNS queries per day

# DNS: a distributed, hierarchical database



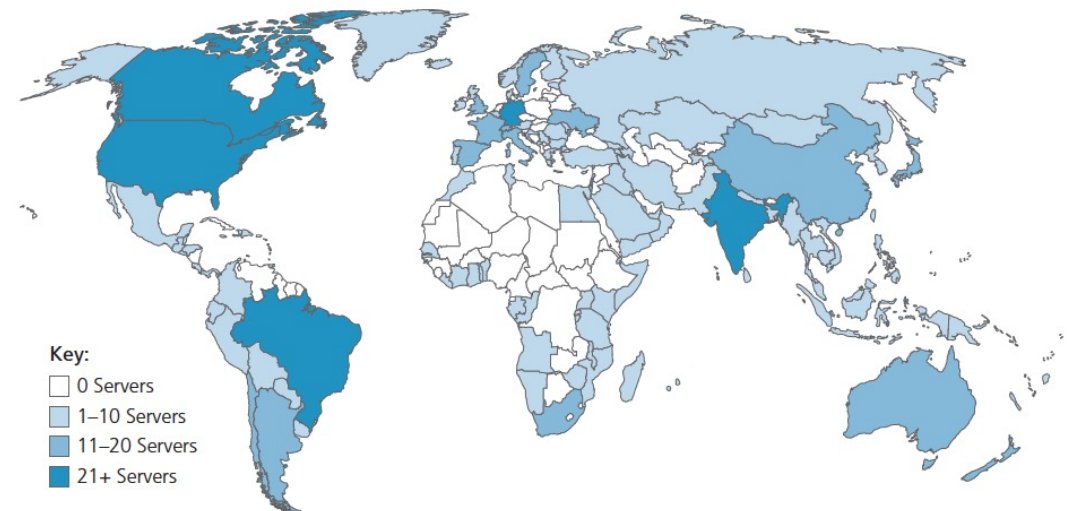
Client wants IP address for [www.amazon.com](http://www.amazon.com); 1<sup>st</sup> approximation:

- client queries root server to find .com DNS server
- client queries .com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for [www.amazon.com](http://www.amazon.com)

# DNS: root name servers

- official, contact-of-last-resort by name servers that cannot resolve name
- *incredibly important* Internet function
  - Internet couldn't function without it!
  - DNSSEC – provides security (authentication and message integrity)
- ICANN (Internet Corporation for Assigned Names and Numbers) manages root DNS domain

Root DNS managed by 13 organizations worldwide, each deploying one logical “server” replicated many times (~200 servers in US, more than 400 worldwide)



# TLD and authoritative servers

## ❑ Top-level domain (TLD) servers:

- ❖ responsible for .com, .org, .net, .edu, .jobs, ...  
and all top-level country domains, e.g.: .uk, .fr, .be
- ❖ Network Solutions maintains servers for .com and .net TLD
- ❖ DNS Belgium for the .be TLD

## ❑ Authoritative DNS servers:

- ❖ organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- ❖ can be maintained by organization or service provider

## Local DNS name server

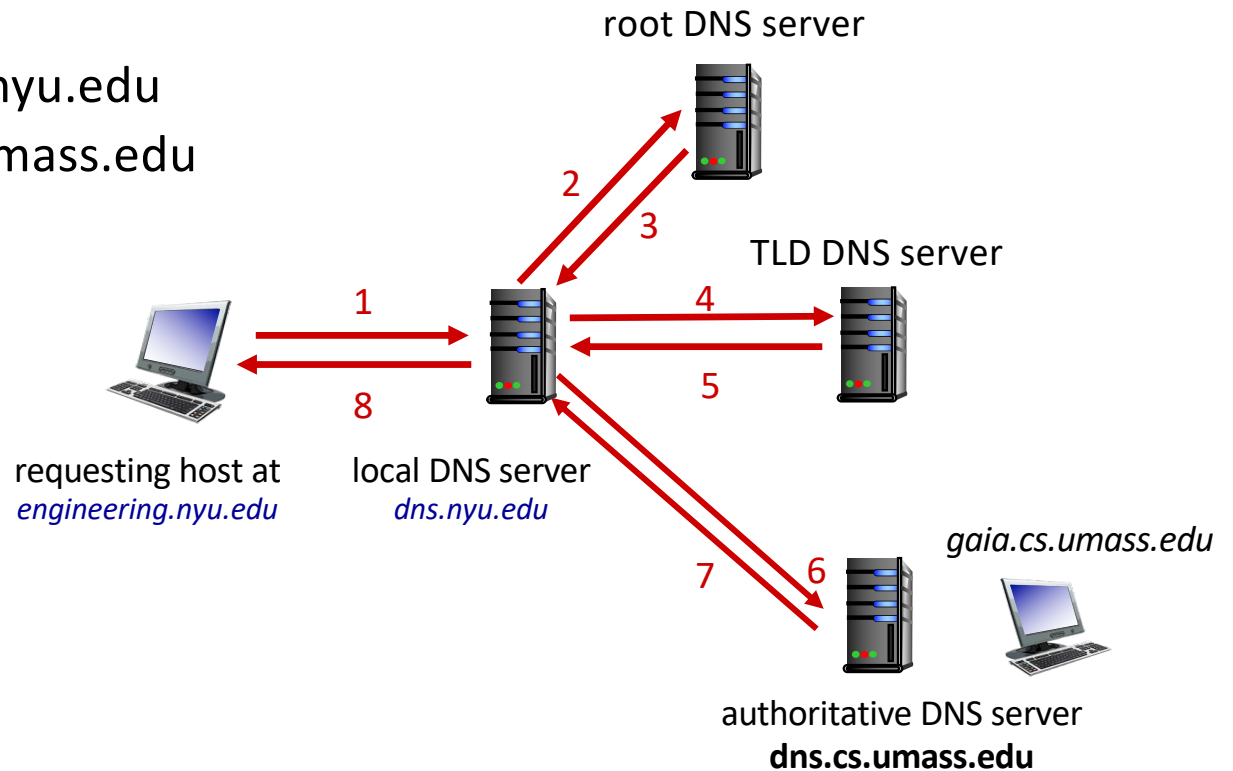
- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one
  - ❖ also called “default name server”
- when host makes DNS query, query is sent to its local DNS server
  - ❖ has local cache of recent name-to-address translation pairs (but may be out of date!)
  - ❖ acts as proxy, forwards query into hierarchy

# DNS name resolution: iterated query

**Example:** host at `engineering.nyu.edu`  
wants IP address for `gaia.cs.umass.edu`

## Iterated query:

- contacted server replies with name of server to contact
- “I don’t know this name, but ask this server”
- caching in local DNS
  - can skip steps



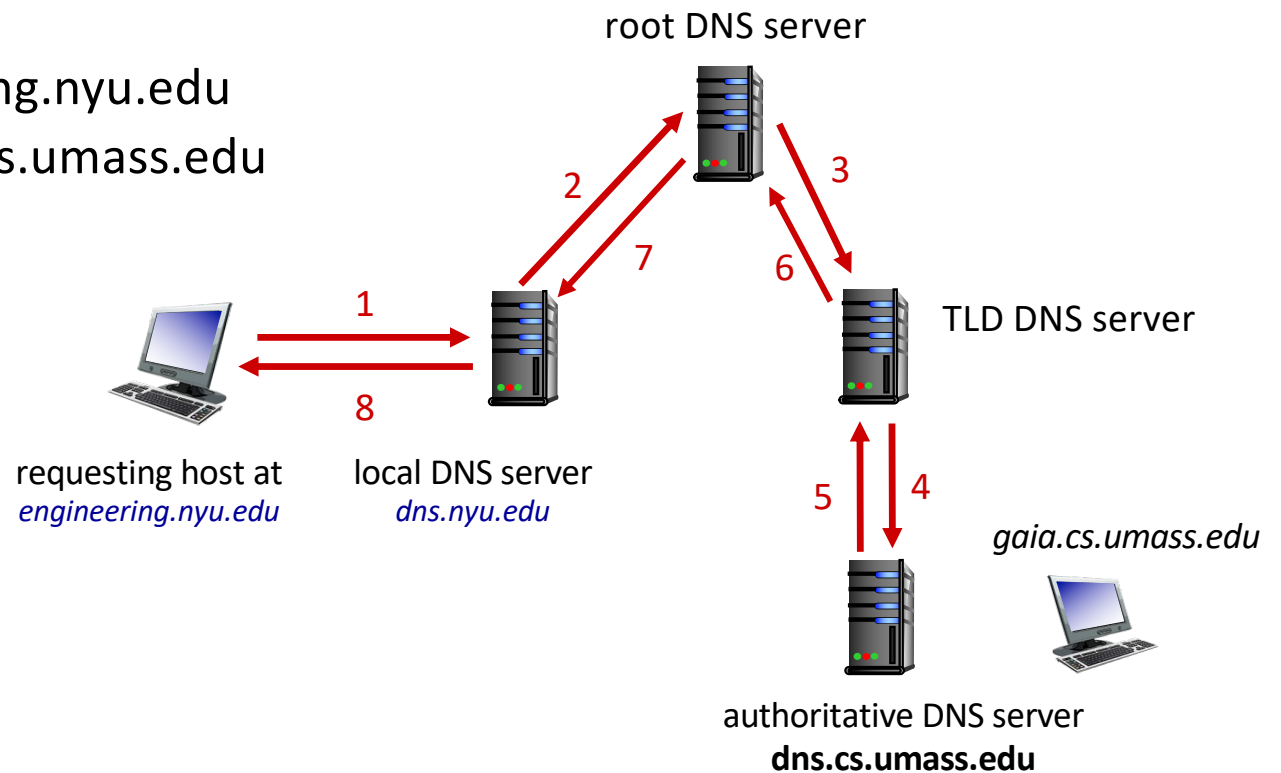


# DNS name resolution: recursive query

**Example:** host at engineering.nyu.edu wants IP address for gaia.cs.umass.edu

## Recursive query:

- puts burden of name resolution on contacted name server
- heavy load at upper levels of hierarchy?
- caching, where? what?



## Caching and updating DNS records

- ❑ once (any) name server learns mapping, it *caches* mapping
  - ❖ cache entries timeout (disappear) after some time (TTL)
  - ❖ TLD servers typically cached in local name servers
    - Thus root name servers not often visited
- ❑ cached entries may be *out-of-date* (best effort name-to-address translation!)
  - ❖ if name host changes IP address, may not be known Internet-wide until all TTLs expire
- ❑ update/notify mechanisms (Dynamic DNS):
  - ❖ RFC 2136

# DNS records

DNS: distributed database storing resource records (RR)

RR format: (name, value, type, ttl)

## □ Type=A (or AAAA)

- ❖ **name** is **hostname** (aka FQDN)
- ❖ **value** is IPv4 (or IPv6) address

## □ Type=NS

- ❖ **name** is **domain**, DNS zone (e.g. amazon.com)
- ❖ **value** is hostname of authoritative name server for this domain
- ❖ domain name ≠ hostname

## □ Type=CNAME

- ❖ **name** is alias name for some “canonical” (the real) name  
www.ibm.com is really  
servereast.backup2.ibm.com
- ❖ **value** is canonical name

## □ Type=MX

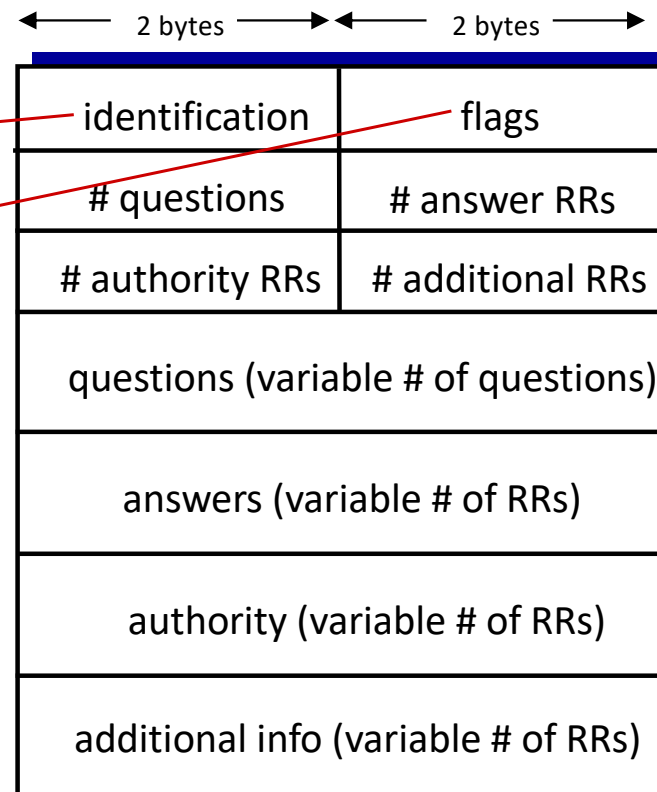
- ❖ **value** is name of mail server associated with the domain **name**

# DNS protocol messages

DNS *query* and *reply* messages, both with same *message format* (usually sent over UDP, why?)

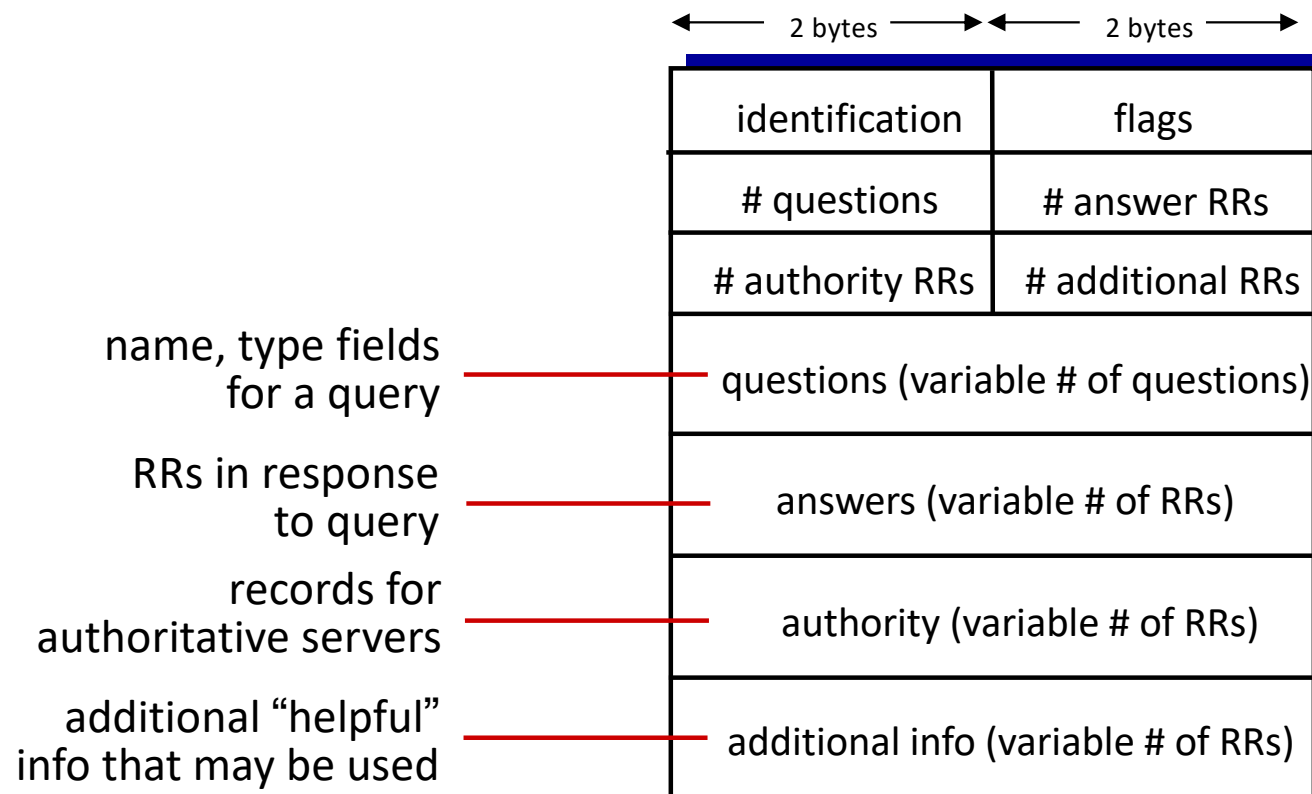
msg header

- ❖ **identification**: 16 bit # for query, reply to query uses same #
- ❖ **flags**:
  - query or reply
  - recursion desired
  - recursion available
  - reply is authoritative



## DNS protocol messages

DNS *query* and *reply* messages, both have same *format*:



# Inserting records into DNS

Example: new startup “Network Utopia”

- register name networkutopia.com at *DNS registrar* (e.g., Network Solutions)
  - provide names, IP addresses of authoritative name server (primary and secondary)
  - registrar inserts two RRs into .com TLD server:
    - (networkutopia.com, dns1.networkutopia.com, NS)
    - (dns1.networkutopia.com, 212.212.212.1, A)
- create authoritative server locally with IP address 212.212.212.1
  - type A record for www.networkutopia.com
  - type MX record for networkutopia.com

## Chapter 2: outline

2.1 Principles of network applications

2.2 Web and HTTP

2.3 The Domain Name System (DNS)

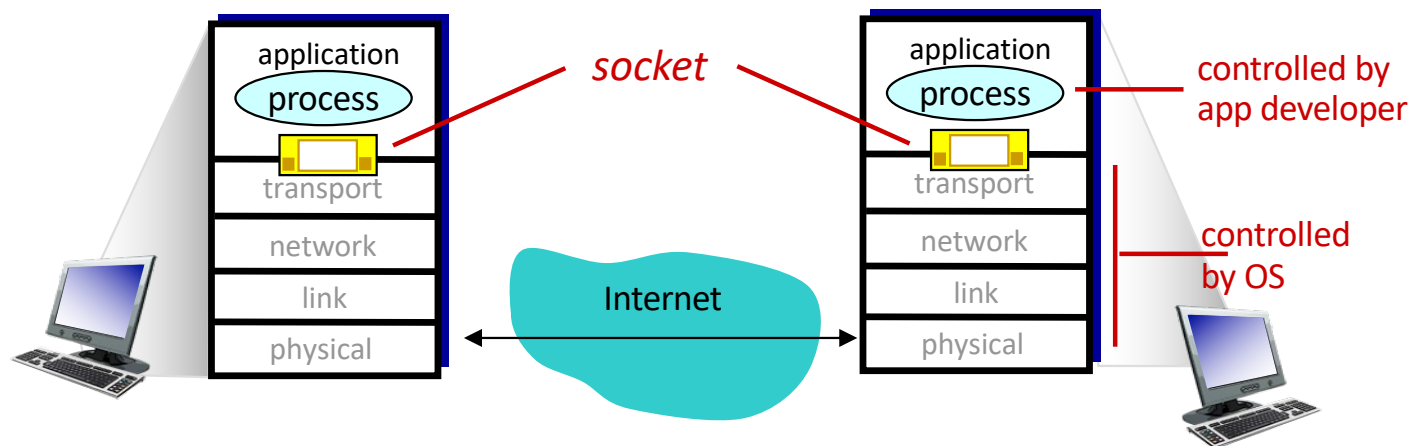
2.4 Socket programming with UDP and TCP

# Socket programming

*goal:* learn how to build client/server applications that communicate using sockets

*socket:* door between application process and end-end-transport protocol

Socket API introduced in BSD4.1 UNIX, 1981





## Socket programming

*Two socket types for two transport services:*

- ❖ *UDP*: unreliable datagram
- ❖ *TCP*: reliable, byte stream-oriented

## Socket programming *with UDP*

UDP: no “connection” between client & server processes

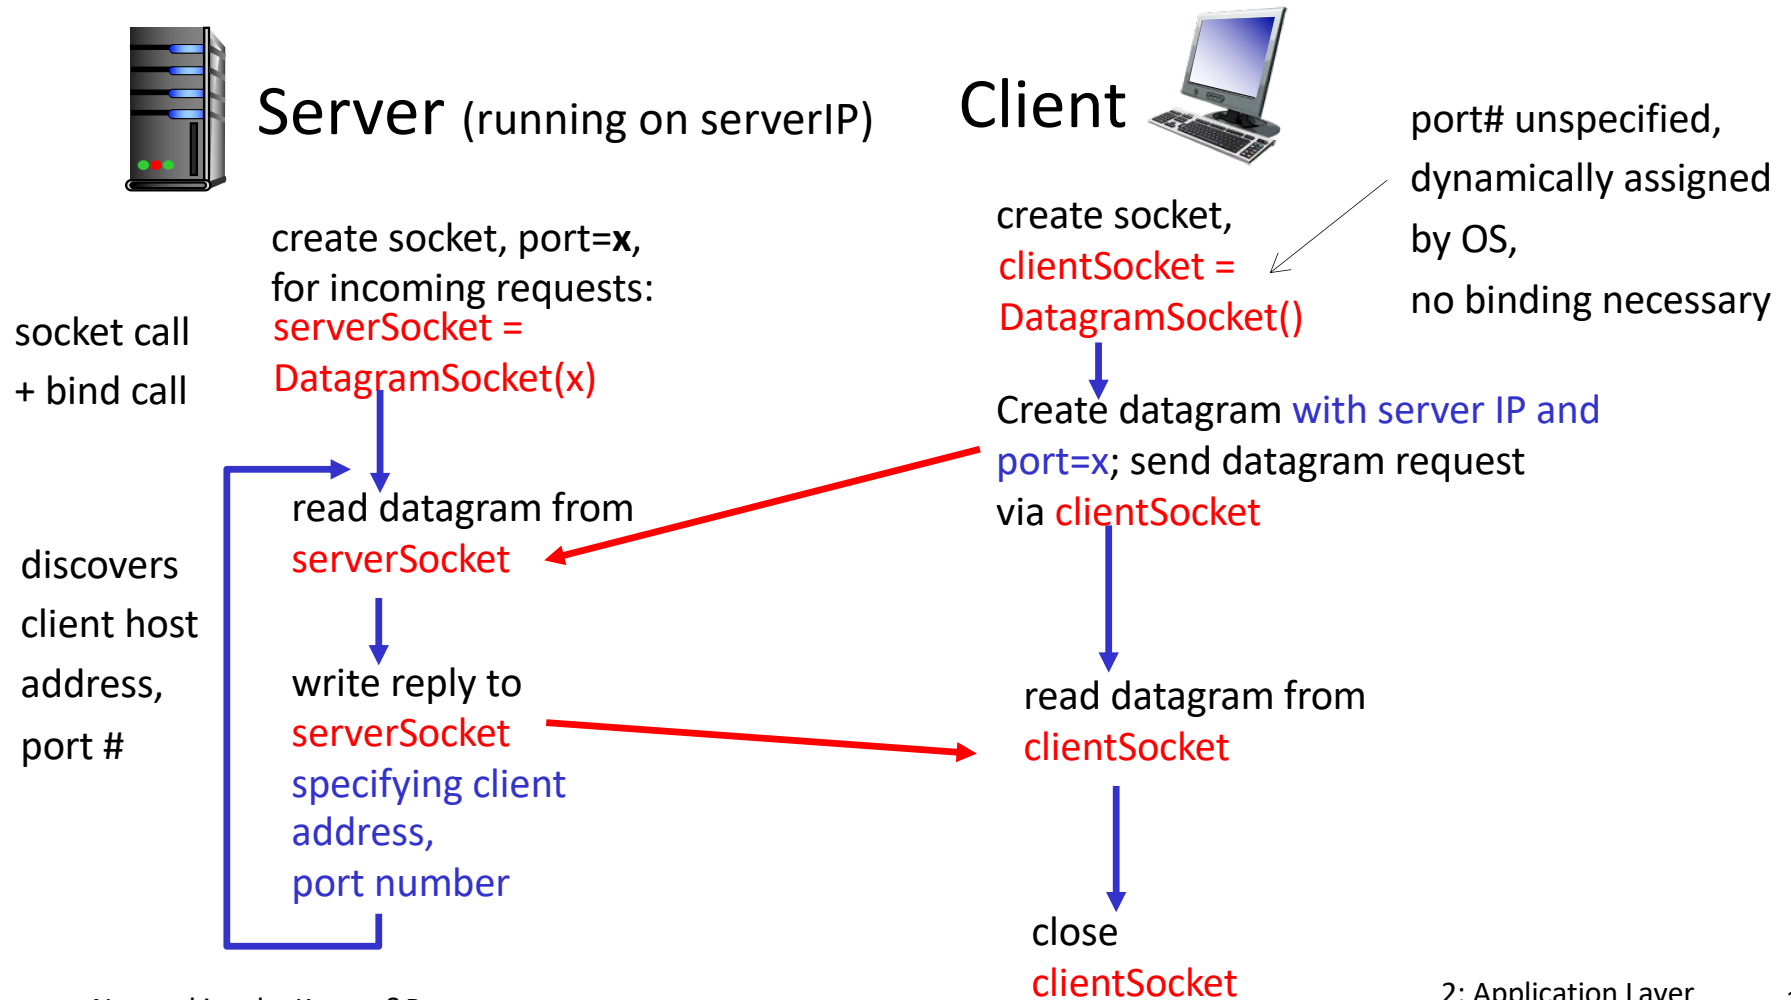
- ❑ no handshaking before sending data
- ❑ sender process explicitly attaches IP destination address and port # to each application data unit (aka datagram)
- ❑ receiver process extracts sender IP address and port# from received datagram

UDP: transmitted datagram may be lost or received out-of-order

Application viewpoint:

UDP provides *unreliable* transfer of datagrams between client and server

# Client/server socket interaction: UDP



# UDP observations & questions

- ❑ Both client and server processes use same DatagramSocket
- ❑ Destination IP address and port # are explicitly attached to datagram by client and server processes when sent through socket
- ❑ Source IP address and port # are discovered by client and server processes when they receive a datagram
- ❑ Can multiple clients use the server?
  - ❖ How many sockets needed by server?
  - ❖ How many port numbers used at server side?

## Socket programming *with TCP*

### Client must contact server

- ❑ server process must first be running
- ❑ server must have created socket (door) that welcomes client's contact (**welcoming socket**)

### Client contacts server by:

- ❑ creating TCP socket specifying IP address, port number of server process
- ❑ When **client creates socket**: client TCP establishes connection to server TCP

- ❑ When contacted by client, **server TCP creates new socket** for server process to communicate with that particular client
  - ❖ allows server to talk with multiple clients
  - ❖ source IP and source port numbers used to distinguish clients (*more in Chap 3*)

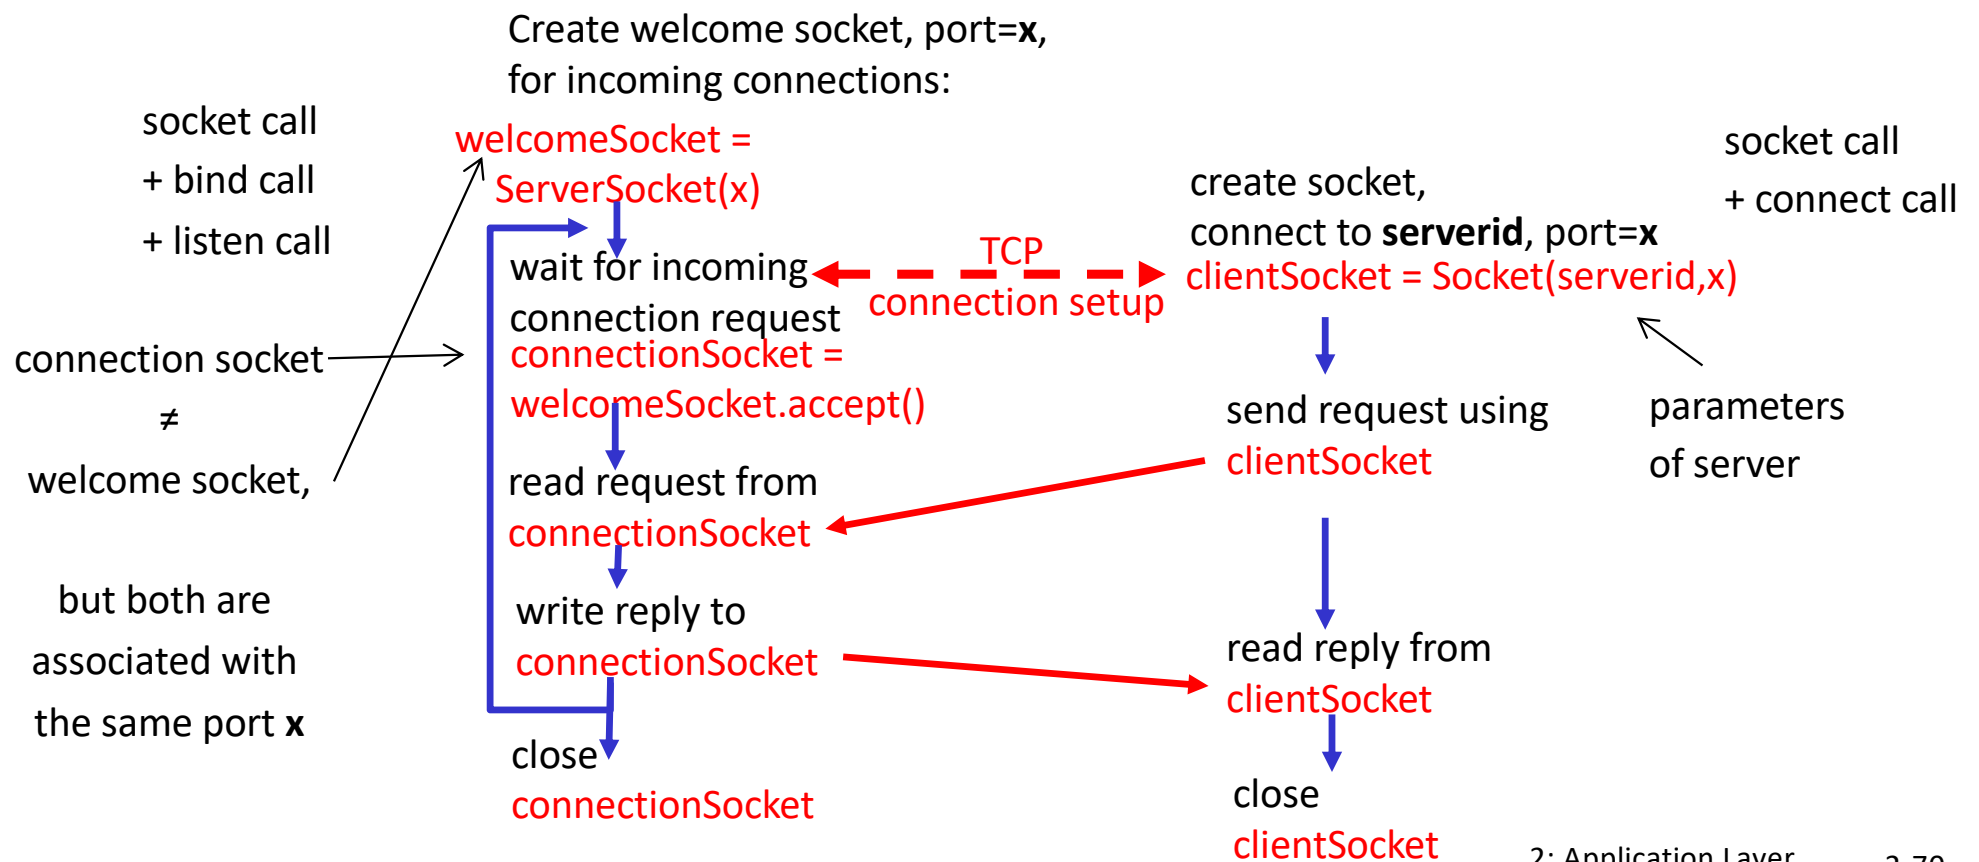
### **application viewpoint**

*TCP provides reliable, in-order byte-stream transfer ("pipe") between client and server*

# Client/server socket interaction: TCP

## Server (running on **serverid**)

## Client



## TCP observations & questions

- ❑ Two types of socket primitives: ServerSocket and Socket
- ❑ When client knocks on server's welcome socket, server creates a dedicated connectionSocket and completes TCP connection
- ❑ Destination IP address and port # are not explicitly attached to application data unit sent and received by client and server processes
  - ❖ The application data unit is just a block of application data bytes
- ❑ Can multiple clients use the server?
  - ❖ How many sockets needed by server?
  - ❖ How many port numbers used at server side?

# Chapter 2: Summary

## □ application architectures

- ❖ client-server
- ❖ P2P

## □ application service requirements:

- ❖ reliability, throughput, delay

## □ Internet transport service model

- ❖ connection-oriented, reliable, byte-stream : TCP
- ❖ unreliable, datagrams: UDP

## □ specific protocols:

- ❖ HTTP
- ❖ DNS

## □ socket programming:

- ❖ TCP, UDP sockets



# Chapter 2: Summary

Most importantly: learned about *protocols*

- typical request/reply message exchange:

- ❖ client requests info or service
- ❖ server responds with data, status code

- message formats:

- ❖ *headers*: fields giving info about data
- ❖ *data* (payload): info being communicated

*Important themes:*

- centralized vs. decentralized
- stateless vs. stateful
- scalability
- reliable vs. unreliable message transfer
- “complexity at network edge”